

# HW3: Paper Reading

Emily Tseng  
et397@cornell.edu

February 27, 2020

## 1 Introduction

In this assignment, we report on the following 3 papers:

1. Lei et al. (2016): *Rationalizing Neural Predictions*
2. 2
3. 3

## 2 Lei et al. (2016): *Rationalizing Neural Predictions*

What is the model? Roughly how many parameters does the model have?

What is the latent structure?

What is the inference approach?

If they use amortized inference, what is the amortized model? How many parameters does it have?

What is the training objective?

What training / inference approximations does the paper make?

Does the paper evaluate interpretability? If so how?

## 3 2

What is the model? Roughly how many parameters does the model have?

What is the latent structure?

What is the inference approach?

If they use amortized inference, what is the amortized model? How many parameters does it have?

What is the training objective?

What training / inference approximations does the paper make?

Does the paper evaluate interpretability? If so how?

## 4 3

What is the model? Roughly how many parameters does the model have?

What is the latent structure?

What is the inference approach?

If they use amortized inference, what is the amortized model? How many parameters does it have?

What is the training objective?

What training / inference approximations does the paper make?

Does the paper evaluate interpretability? If so how?

## 5 Discussion

## 6 Conclusion

## References

Lei, T., Barzilay, R., and Jaakkola, T. (2016). Rationalizing neural predictions. *arXiv preprint arXiv:1606.04155*.